



Kubernetes Interview Questions & Troubleshooting Answers



A Practical Handbook for DevOps
Job Seekers



1. What happens when a pod is stuck in ContainerCreating?

Purpose: To evaluate your knowledge of the pod lifecycle and your ability to investigate container startup problems.

Context:

This usually occurs when the scheduler has assigned a node, but the container runtime cannot finish setting up the container. Common causes include image pull issues, missing secrets/configs, or volume mount delays.

Troubleshooting Answer:

1. `kubectl describe pod <pod-name>` # Check events
2. `kubectl get pv,pvc` # Check volume bindings
3. `kubectl logs <pod-name>` # See if image pull is the issue

Look for: ImagePullBackOff, volume mount issues, network errors.



2. How do you handle a CrashLoopBackOff error?

Purpose: To assess your troubleshooting ability for runtime and application configuration failures.

Context:

Indicates the pod started, crashed due to an application issue, and is now in a restart loop. Often seen with misconfigured environment variables or bad startup commands.

Troubleshooting Answer:

1. `kubectl logs <pod-name> --previous`
2. `kubectl describe pod <pod-name>`
3. `kubectl get configmap,secret`

Look for: Misconfigured probes, environment issues, missing dependencies.



3. A node is in NotReady state. What steps do you take?

Purpose: To confirm your understanding of cluster health and how to diagnose node-level issues.

Context:

Nodes go NotReady when the kubelet is not reporting healthy status to the API server, often due to resource pressure or network failures.

Troubleshooting Answer:

1. `kubectl get nodes`
2. `ssh <node> "sudo systemctl status kubelet"`
3. `df -h` # Check disk space

Look for: Kubelet failures, disk/memory saturation, network unreachable.



4. DNS resolution is failing inside a pod. How will you fix it?

Purpose: To verify your knowledge of internal DNS resolution and network policy behavior within the cluster.

Context:

Internal or external names can't be resolved inside pods. This usually relates to CoreDNS failure, a broken **resolv.conf**, or network policy issues.

Troubleshooting Answer:

1. `kubectl exec -it <pod> -- nslookup google.com`
2. `kubectl get pods -n kube-system | grep coredns`
3. `kubectl logs <coredns-pod> -n kube-system`

Look for: CoreDNS pod failures, incorrect network policies.



5. Your service is not reachable via ClusterIP. Why?

Purpose: To check your understanding of Kubernetes service discovery and endpoint bindings.

Context:

A service created with ClusterIP is not routing traffic to backend pods. Often due to incorrect selectors or port mismatches between service and pod.

Troubleshooting Answer:

1. `kubectl describe svc <svc-name>`
2. `kubectl get endpoints <svc-name>`
3. `kubectl get pods -l app=<label>`

Look for: Label mismatches, port mismatches, missing endpoints.



6. Why would a PVC be stuck in Pending?

Purpose: To evaluate your understanding of persistent storage provisioning and storage class configuration.

Context:

A PersistentVolumeClaim can't bind to a suitable PersistentVolume due to storage class mismatches or no available capacity.

Troubleshooting Answer:

1. `kubectl get pvc`
2. `kubectl describe pvc <pvc-name>`
3. `kubectl get sc` # Check storage class

Look for: No matching PV, misconfigured storage class.



7. kubectl exec and logs are hanging. How do you troubleshoot?

Purpose: To confirm your ability to diagnose communication problems between control plane and nodes.

Context:

This typically indicates networking issues between the control plane and the kubelet, or the container runtime is unresponsive.

Troubleshooting Answer:

1. `kubectl describe pod <pod-name>`
2. `kubectl get pod -o wide # Get IP`
3. `kubectl top node # Check resource pressure`

Look for: Network issues, kubelet health, resource bottlenecks.



8. Ingress is not routing traffic. What do you do?

Purpose: To assess your understanding of traffic routing, Ingress configuration, and backend service connectivity.

Context:

Ingress is configured, but requests are not reaching the backend. May be caused by incorrect host/path settings or broken Ingress controller setup.

Troubleshooting Answer:

1. `kubectl describe ingress <name>`
2. `kubectl logs <ingress-controller-pod>`
3. `kubectl get svc -n ingress-nginx`

Look for: DNS mismatch, backend unreachable, wrong path/host rules.



9. Why is HPA not scaling your pods?

Purpose: To determine your awareness of autoscaling requirements and monitoring setup in Kubernetes.

Context:

HPA relies on metrics-server to retrieve CPU/memory usage. If metrics are missing or below thresholds, no scaling occurs.

Troubleshooting Answer:

1. `kubectl get hpa`
2. `kubectl top pods`
3. `kubectl get pods -n kube-system | grep metrics-server`

Look for: No metrics-server, thresholds not met, misconfigured CPU/memory targets.



10. You updated a Deployment but pods didn't change. Why?

Purpose: To test your ability to manage deployment strategies and version control for applications.

Context:

Often happens when the manifest has no actual spec change or uses an unchanged image tag. Kubernetes sees no diff and doesn't roll out.

Troubleshooting Answer:

1. `kubectl rollout status deployment <name>`
2. `kubectl rollout history deployment <name>`
3. `kubectl describe deployment <name>`

Look for: No actual spec change, same image tag used, caching.



Bavithran
@bavicnative



Found this useful?



Follow



Let's connect



Found it useful? Drop your thoughts below and share it with your fellow DevOps engineers!

Follow me for more DevOps, Kubernetes, and cloud-native insights.



Repost